Jaywalk! for all I care...

Introduction: Automated driver assistance systems (ADAS) face challenges with pedestrians, especially in complex urban environments where pedestrians exhibit diverse and unpredictable behaviors. The uncertainty exhibited by these road users - in addition to the can-of-worms in collecting it (Feng et al., 2021) - is magnified in automated driving systems (ADS) designed to execute both the dynamic driving task (DDT) and strategic tasks (i.e., trip planning, route selection). However, this off-loading of the driving task requires an adequately scaled system that can handle a myriad of ambiguous pedestrian behaviors (Mitman and Ragland, 2007). Herein lies the gap between the amount of training possible to provide such systems and the training necessary to ensure *some* level of safe deployment in multimodal contexts.

Background: Reinforcement learning (RL) has gained significant popularity on this front for its flexibility to explore a state space independently, providing resiliency towards unseen events and environments. Reasonably, a lack of interpretability becomes concerning in high-risk domains, where agents act according to learned optimal policies not entirely understood (Armstrong and Mindermann, 2018; Glanois et al., 2021). Mistrust leads to distrust: before real-world applications of RL models are statistically real test-beds for them. However, creating such simulations with distribution-level accuracy is difficult for sparse events struggling with dimensionality (i.e., safety-critical driving events, Yan et al. (2023)). At scale, such is also the case for pedestrians, whom inherently and indefinitely add complexity to their environments through interaction (autocirricula, Leibo et al. (2019)). But fear not! Wang et al. (2022a,b, 2020), Wang et al. (2016, 2021), ..., have made significant efforts developing human-like driving models to exhibit socially compatible behaviors. In light, socially compatible behaviors for various kinds of road agents must be learned (i.e., vulnerable road users (pedestrians) require more sensitivity in interaction terms of space and time).

Research question: How do we walk? Due to collection costs and privacy concerns, robust (and granular) pedestrian behavior data is scarce, which limits the amount of knowledge we have about pedestrian path planning at a microenvironment scale. It may be interesting to use observed expert demonstrations to explore, for example, how pedestrians avoid re-routes due to construction (flexibility to new tasks), or what running to the bus looks like (robustness to goal specification). Johora and Müller (2021) suggests pedestrians often path plan irrationally, is it preferential? Because there are multiple expert demonstrations per participant, it would also be interesting to explore preferences among pedestrian in urban navigation.

Analysis plan: Let's explore walking. Maybe potentially with deep imitation learning and expert demonstration. These demonstrations are provided via previously collected GIS-linked walking activity from pedestrians in Seattle, WA. Of note, this dataset contains multiple traversals for ~ 300 pedestrians.

I would be remiss to not mention some recent work. What initially sparked interest into expert demonstration for intelligent planning was Rhinehart et al. (2018), who model continuous-state, discrete-time, partially-observed Markov processes in a deep imitative learning model (an exact target no longer, cars can avoid pot-holes). Vozniak et al. (2020) uses deep imitation learning to explore pedestrian behaviors in critical safety scenarios, and Trinh et al. (2020) employed a PPO-based path planning model that followed common walking conventions and human behaviours such as walking on the left side of the road and staying away from dangerous obstacles. Honorable mention to a devil's advocate, Lake et al. (2017), who kinda thinks the whole things a joke.

Next steps: Will need help scaling scope/focus. Get feedback from EV, meet to discuss scope and brainstorm implementation. Please reach out if you have significant concerns!

References

Armstrong, S., Mindermann, S., 2018. Occam's razor is insufficient to infer the preferences of irrational agents. Adv. Neural Inf. Process. Syst. 31. URL: https://proceedings.neurips.cc/paper/2018/file/d89a66c7c80a29b1bdbab0f2a1a94af8-Paper.pdf.

Feng, Y., Duives, D., Daamen, W., Hoogendoorn, S., 2021. Data collection methods for studying pedestrian behaviour: A systematic review. Build. Environ. 187, 107329. URL: https://www.sciencedirect.com/science/article/pii/S0360132320306983, doi:10.1016/j.buildenv.2020.107329.

Glanois, C., Weng, P., Zimmer, M., Li, D., Yang, T., Hao, J., Liu, W., 2021. A survey on interpretable reinforcement learning URL: http://arxiv.org/abs/2112.13112, arXiv:2112.13112.

Johora, F.T., Müller, J.P., 2021. Modeling interactions of multimodal road users in shared spaces URL: http://arxiv.org/abs/2107.02083, arXiv:2107.02083.

Lake, B.M., Ullman, T.D., Tenenbaum, J.B., Gershman, S.J., 2017. Building machines that learn and think like people. Behav. Brain Sci. 40, e253. URL: http://dx.doi.org/10.1017/S0140525X16001837, doi:10.1017/S0140525X16001837.

Leibo, J.Z., Hughes, E., Lanctot, M., Graepel, T., 2019. Autocurricula and the emergence of innovation from social interaction: A manifesto for Multi-Agent intelligence research URL: http://arxiv.org/abs/1903.00742, arXiv:1903.00742.

Mitman, M.F., Ragland, D.R., 2007. Crosswalk confusion: More evidence why pedestrian and driver knowledge of the vehicle code should not be assumed. Transp. Res. Rec. 2002, 55--63. URL: https://doi.org/10.3141/2002-07, doi:10.3141/2002-07.

Rhinehart, N., McAllister, R., Levine, S., 2018. Deep imitative models for flexible inference, planning, and control URL: http://arxiv.org/abs/1810.06544, arXiv:1810.06544.

Trinh, T.T., Vu, D.M., Kimura, M., 2020. A pedestrian path-planning model in accordance with obstacle's danger with reinforcement learning, in: Proceedings of the 3rd International Conference on Information Science and Systems, Association for Computing Machinery, New York, NY, USA. pp. 115--120. URL: https://doi.org/10.1145/3388176.3388187, doi:10.1145/3388176.3388187.

Vozniak, I., Klusch, M., Antakli, A., Müller, C.A., 2020. InfoSalGAIL: Visual attention-empowered imitation learning of pedestrian behavior in critical traffic scenarios. Int Jt Conf Comput Intell , 325--337URL: https://pdfs.semanticscholar.org/f10e/c4e8446dc9bc4b5d8b88b6a0b74f4207af03.pdf, doi:10.5220/0010020003250337.

Wang, L., Sun, L., Tomizuka, M., Zhan, W., 2020. Socially-Compatible behavior design of autonomous vehicles with verification on real human data URL: http://arxiv.org/abs/2010.14712, arXiv:2010.14712.

Wang, Q., Chen, J., Ma, J., 2021. Experimental study on individual level interaction between bicycle and pedestrian. J. Stat. Mech. 2021, 093403. URL: https://iopscience.iop.org/article/10.1088/1742-5468/ac1d54/meta, doi:10.1088/1742-5468/ac1d54.

Wang, S., Gwizdka, J., Chaovalitwongse, W.A., 2016. Using wireless EEG signals to assess memory workload in the n-Back task. IEEE Transactions on Human-Machine Systems 46, 424--435. URL: http://dx.doi.org/10.1109/THMS.2015.2476818, doi:10.1109/THMS.2015.2476818.

Wang, W., Wang, L., Zhang, C., Liu, C., Sun, L., 2022a. Social interactions for autonomous driving: A review and perspectives URL: http://arxiv.org/abs/2208.07541, arXiv:2208.07541.

Wang, W., Wang, L., Zhang, C., Liu, C., Sun, L., 2022b. Social interactions for autonomous driving: A review and perspectives URL: http://arxiv.org/abs/2208.07541, doi:10.1561/2300000078, arXiv:2208.07541.

Yan, X., Zou, Z., Feng, S., Zhu, H., Sun, H., Liu, H.X., 2023. Learning naturalistic driving environment with statistical realism. Nat. Commun. 14, 2037. URL: http://dx.doi.org/10.1038/s41467-023-37677-5, doi:10.1038/s41467-023-37677-5.